

# Benchmarking Vector Databases on Code Embeddings

---

CS 854 – Performance Engineering  
Fall 2023, Ali Mashtizadeh

Vikram **Subramanian**, Raymond **Chang**, Yahya **Jabary**

*N-dimensional, float*

# Benchmarking Vector Databases on Code Embeddings


*Vector representations of documents*

# What are vector databases?

- Specifically designed to store and retrieve vectors
- Given query vector  $\mathbf{q}$  **[0, 0.3, 0.1]**
  - Retrieves **approximately** similar vectors to  $\mathbf{q}$
  - Similarity depends on DB **load**
- Ex: MilvusDB, ChromaDB

# How do they work? - Similarity


## Cosine Similarity


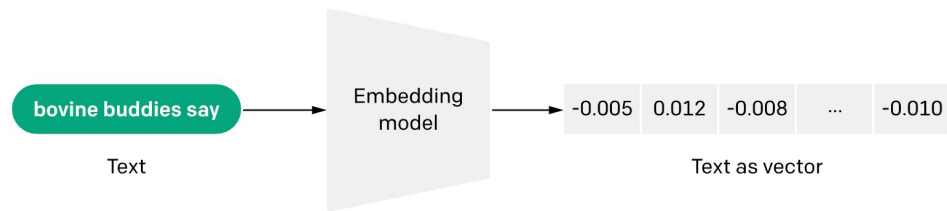


$$\text{sim}(\mathbf{a}, \mathbf{b}) = \frac{\mathbf{a} \cdot \mathbf{b}}{||\mathbf{a}|| \cdot ||\mathbf{b}||}$$

## How do they work? - Indexing


# Navigable Small Worlds





# Why should I care?

- **Novelty:**
  - No standard test dataset for code
  - New Vector DBs every month
- **Practicality:** used in many AI/ML apps
  - Searching and clustering files of code by semantics



# What did we do?

## Contributions:

- A very large high-dimension (1024 dimension) dataset
- First code benchmarking dataset
- Some basic benchmarking



# Why aren't there are many large high-dimensional datasets?

- Computationally expensive
  - $5 \cdot 10^5$  functions in our dataset and we use an embedding of 1024
  - We have  $3 \cdot 10^9$  floats- We need to do  $3 \cdot 10^{18}$  operations
  - Will take 1000 H100s, 3 days
- Solution: Precompute nearest neighbors by generating mutations of code blocks

# Step 1)

Getting initial dataset

# Github API

**Goal:** 500.000+ files of code


## Github API rate limits:

- a. anonymous: 60 requests/h → +8,333h
- b. authenticated: 5000 requests/h → +100h
- c. enterprise: 15000 requests/h → +33h (and lots of money)

# How does rate limiting work?

By identifying you through your:

- IP address, geolocation (default)
- API auth key
- User session cookies (mostly Akamai)
- HTTP requests (mostly Cloudflare)
- TLS fingerprints (very rare)



# Rotating IPs via XFF vulnerability

Rotating proxies for each request through:

- a. AWS API Gateway's large IP pool
- b. BurpSuite PortSwigger IProtate

## HTTP Request:

```
GET /api/v1/otp/check HTTP/1.1
Host: vulnerable-website.com
Content-Type: application/x-www-form-urlencoded
Content-Length: 9
X-Forwarded-For: <client>, <proxy1>, <proxy2>
```

**Result: Unsuccessful**

- Github detects popular proxy pool headers like “X-Amzn-Trace-Id”
- We don't want to pay for custom IP pools

# Fake resource paths

No semantics: %00, %09, %0a, %0c, %20, ...

Bare endpoint: /api/v4/endpoint

/api/v4/endpoint

/api/v4/Endpoint

/api/v4/EndPoint

/api/v4/endpoint%00

/api/v4/%0aendpoint

/api/v4/endpoint%09

/api/v4/%20endpoint

Fake query params:

/api/v4/%20endpoint?phone=+17342239011&code[ ]=123456&code[ ]=654321&...&code[ ]=331337

Result: **Unsuccessful**

# Token rotation / Overwriting ETags

```
$ curl -i https://api.github.com/users/defunkt
> HTTP/2 200
> ...
> etag: W/"61e9...f26f0"

$ curl -i -H 'If-None-Match: "61e9...f26f0"' \
https://api.github.com/users/sueszli
> HTTP/2 304
```

Result: **Unsuccessful**

# Bandwidth throttling

```
from ratelimit import limits, RateLimitException
from backoff import on_exception, expo

import requests

@on_exception(expo, RateLimitException, max_tries=8)
@limits(calls=15, period=900)                                # <-- limit requests
def call_api(url):
    response = requests.get(url)

    if response.status_code != 200:
        raise Exception('API response: {}'.format(response.status_code))
    return response
```

**Result: Unsuccessful** – now we're **even slower**



# Scraping the Github “search” page

- Lets you query by language **and repository size**
- Only displays 5 pages when you get rate limited
- Scraped around 500 entries manually

Result: **Unsuccessful**

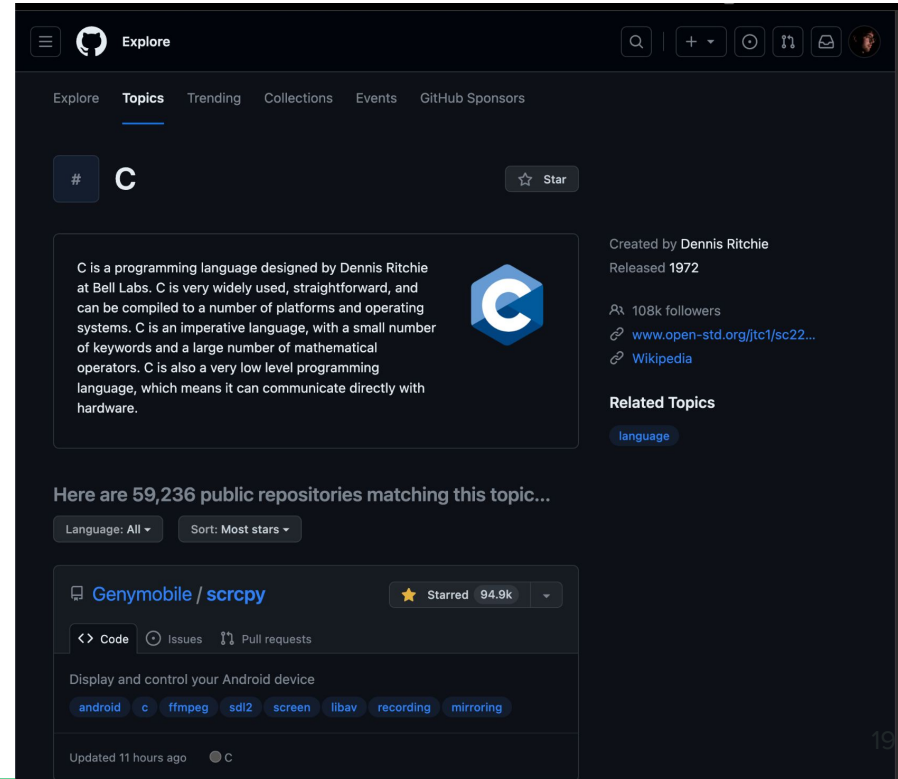


Github, why?

# Scraping the Github “topics” page

- Novel discovery:  
Single page scraping
- 2000 repositories per page  
or ~15 repos per second

Result: **Very Successful**



# Final steps

1. Scrape links to repositories through “Github topics”  
Total repositories: 3,494
2. Clone
3. Filter out non-code  
Total files: 649,257
4. Chunk into 1,000 file commits (>600 commits in ~2 days)  
Otherwise you have to pay for Github LFS

## **Step 2)**

Mutating data to  
generate clusters

# Mutation

- Must generate semantically similar code blocks
- Mutation:
  - Insert block of dead code
- For each function generate five mutated versions
  - Mutations increase in size

# Mutation Example

Original

```
def foo():  
    ...
```

Mutations

```
1. def foo():  
    if False:  
        i = 0  
    ...
```

```
2. def foo():  
    if False:  
        i = 0  
        while(false):  
            ...  
    ...
```



JSON

# Mutation analysis

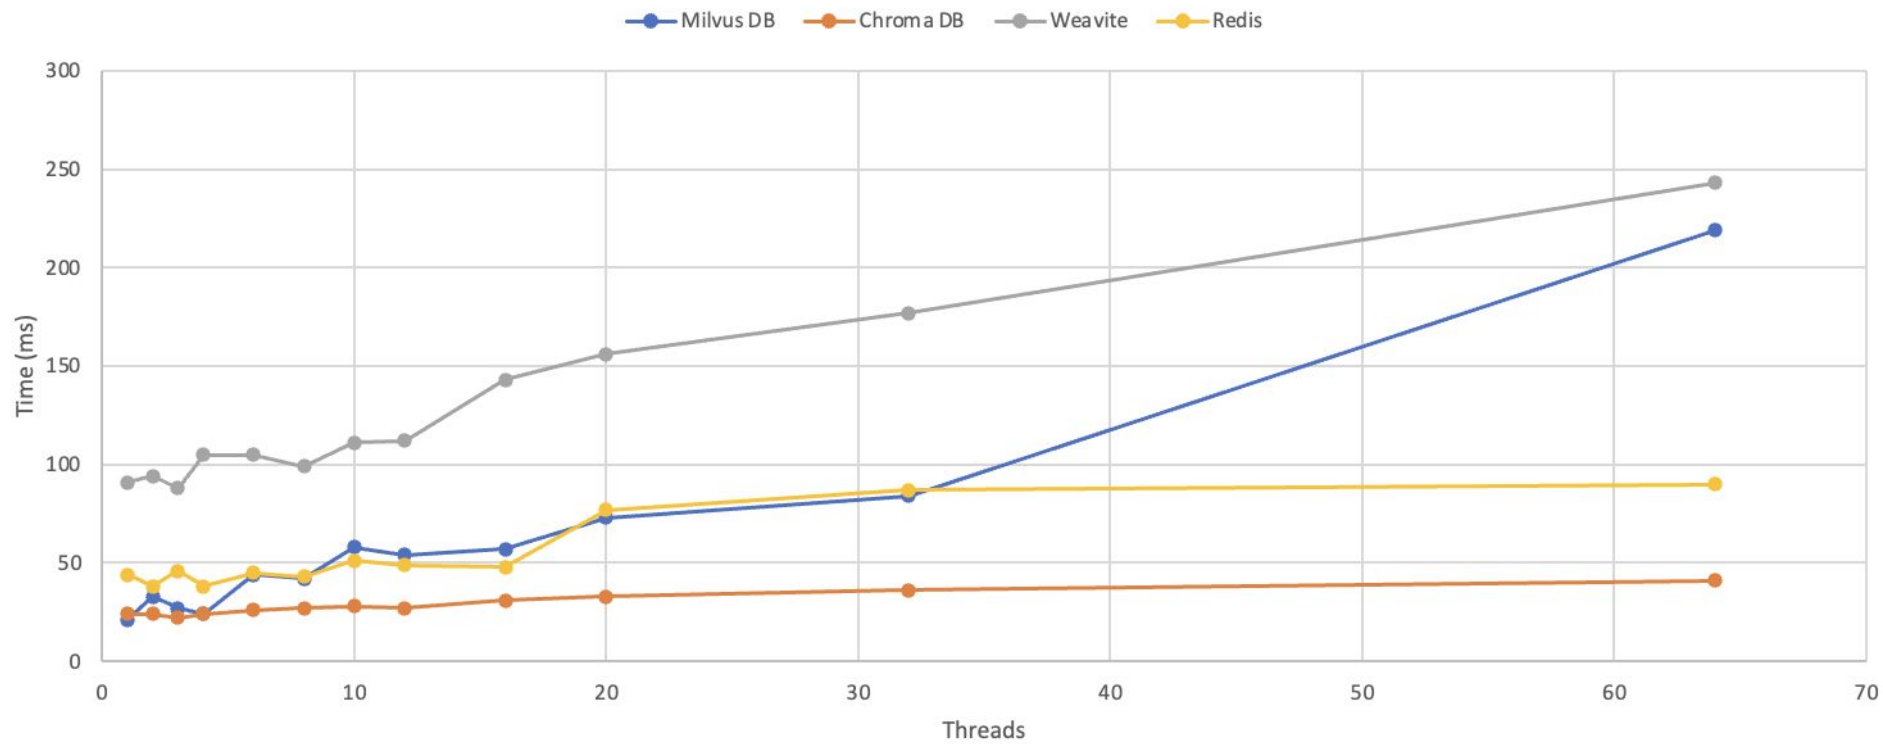
- Want Recall vs db load
- Must ensure mutations actually decrease similarity
- Compare
  - Similarity of the two largest
    - 0.68
  - Smallest function: Similarity with mutations
    - 0.82 to 0.92
  - Largest function: Similarity with mutations
    - 0.85 to 0.91



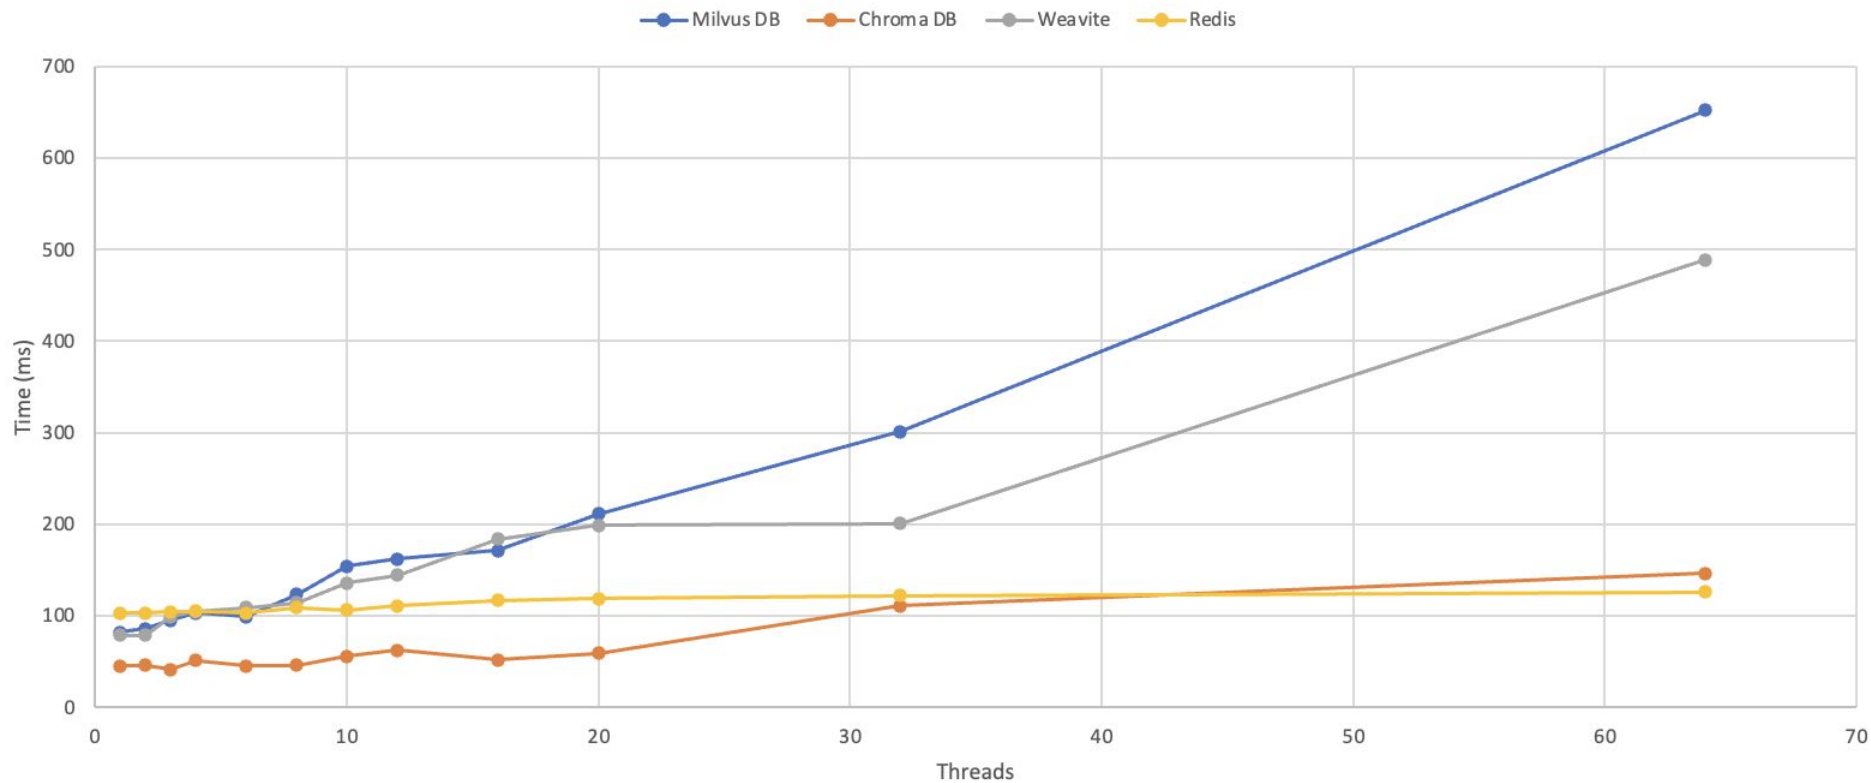
# Step 3)

## Benchmarks

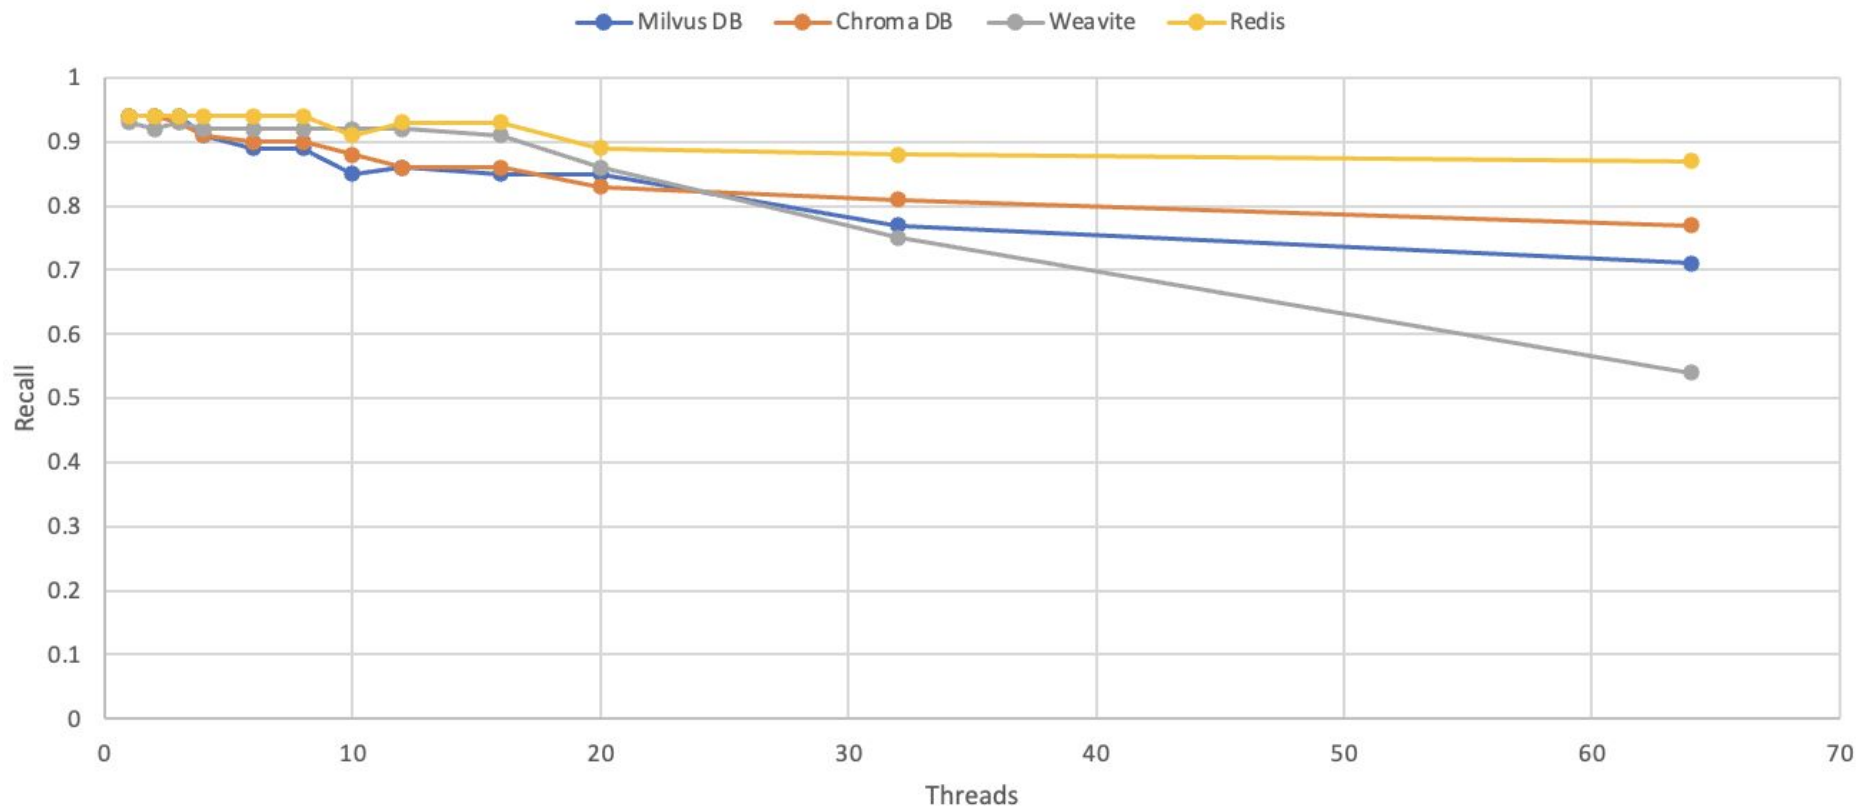
## Write Speeds



## Read speeds



## Recall



# Threats to validity of this approach

Even on a perfect system, recall  $< 1$ .

I.e., our data is imperfect

I.e., given a code block  $x$  and mutation  $x'$ , there exists a code block  $y$  whose cosine similarity to  $x$  is greater than the cosine similarity between  $x$  and  $x'$

Databases under minimal load produced a recall of 0.99 with a perfect dataset. The same DB produced a recall of 0.94.

# Conclusion

- Created a large high-dimensional dataset for VectorDB benchmarking
- Used a novel method to create a code dataset
- Benchmarked some popular vector databases

milvus-io/pymilvus

#1806 **[Bug]: Type issue in orm/schema.py - Incorrect type for X an...**



0 comments




**vikramsubramanian** opened on December 2, 2023



**Thanks for listening!**

# How do they work? - Indexing

## Hierarchical Navigable Small Worlds





# How do they work? - Indexing

## Hierarchical Navigable Small Worlds

